

Python FUSE

File-System in USErspace

Beyond the Traditional File-Systems

<http://mbertozzi.develer.com/python-fuse>

Talk Overview

- ⦿ What is a File-System
- ⦿ Brief File-Systems History
- ⦿ What is FUSE
- ⦿ Beyond the Traditional File-System
- ⦿ API Overview
- ⦿ Examples (Finally some code!!!)
<http://mbertozzi.develer.com/python-fuse>
- ⦿ Q&A

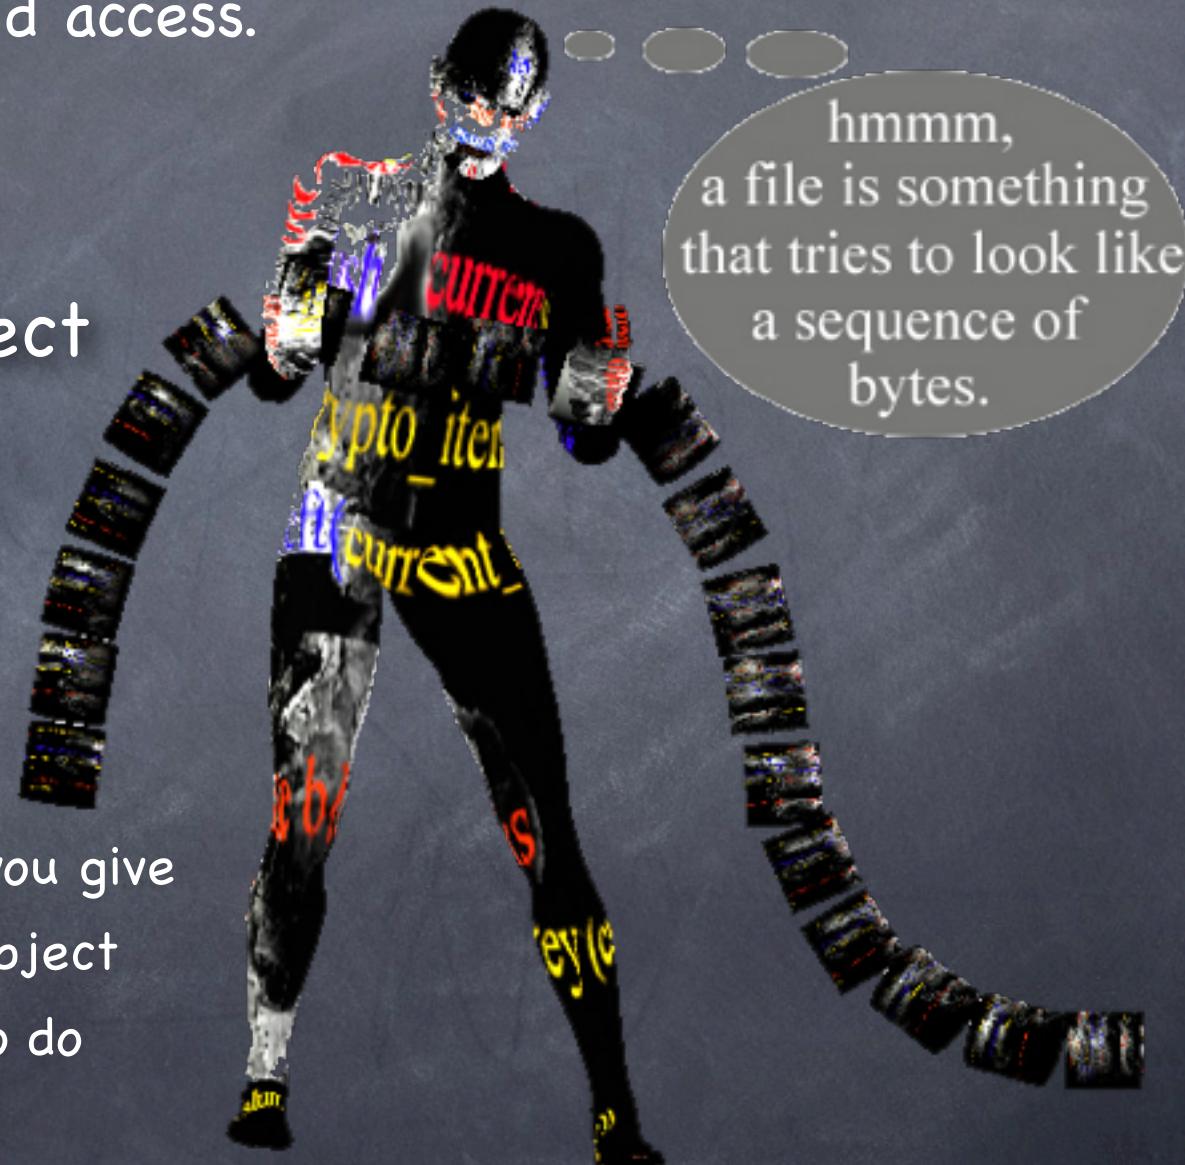
What is a File-System

Is a Method of storing and organizing data
to make it easy to find and access.

...to interact with an object

You name it, and you say
what you want it do.

- The Filesystem takes the name you give
- Looks through disk to find the object
- Gives the object your request to do something.

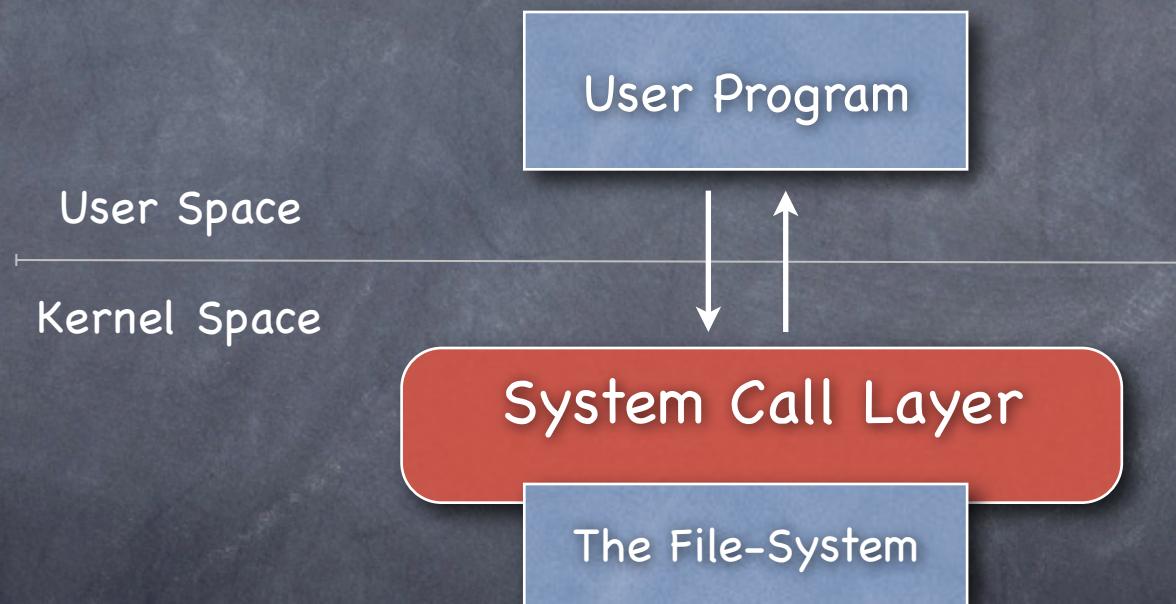


What is a File-System

- ⦿ On Disk Format (...serialized struct)
ext2, ext3, reiserfs, btrfs...
- ⦿ Namespace
(Mapping between name and content)
/home/th30z/, /usr/local/share/test.c, ...
- ⦿ Runtime Service: open(), read(), write(), ...

...A bit of History

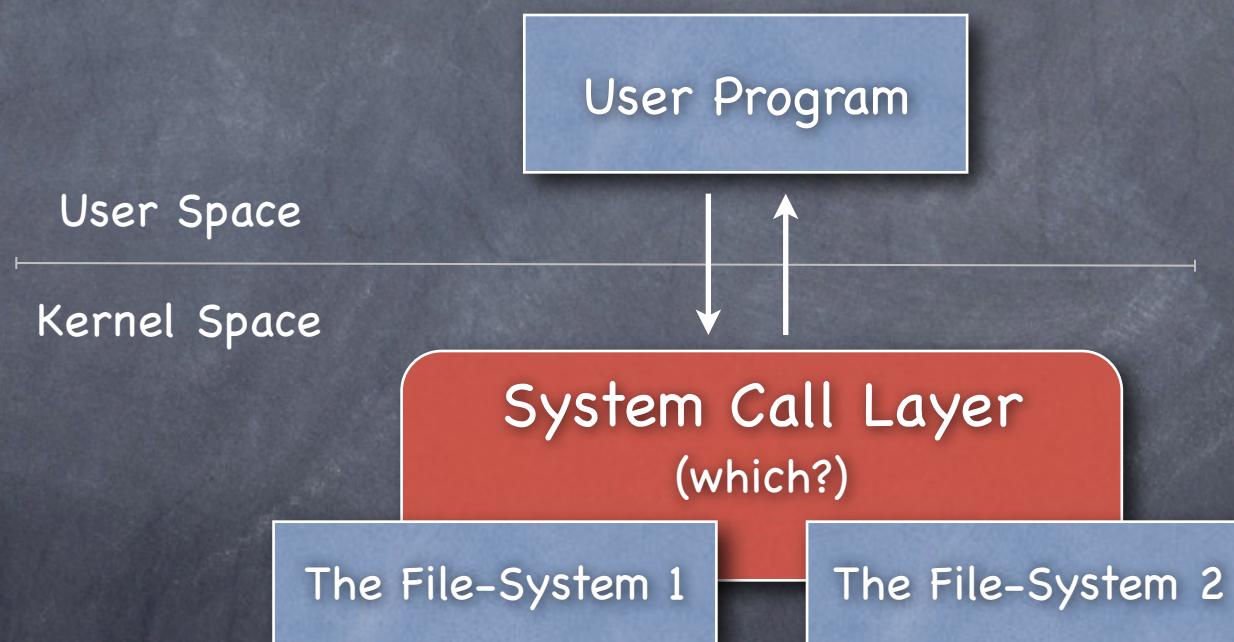
- Multics 1965 (File-System Paper)
A General-Purpose File System For Secondary Storage
- Unix Late 1969



Only One File-System

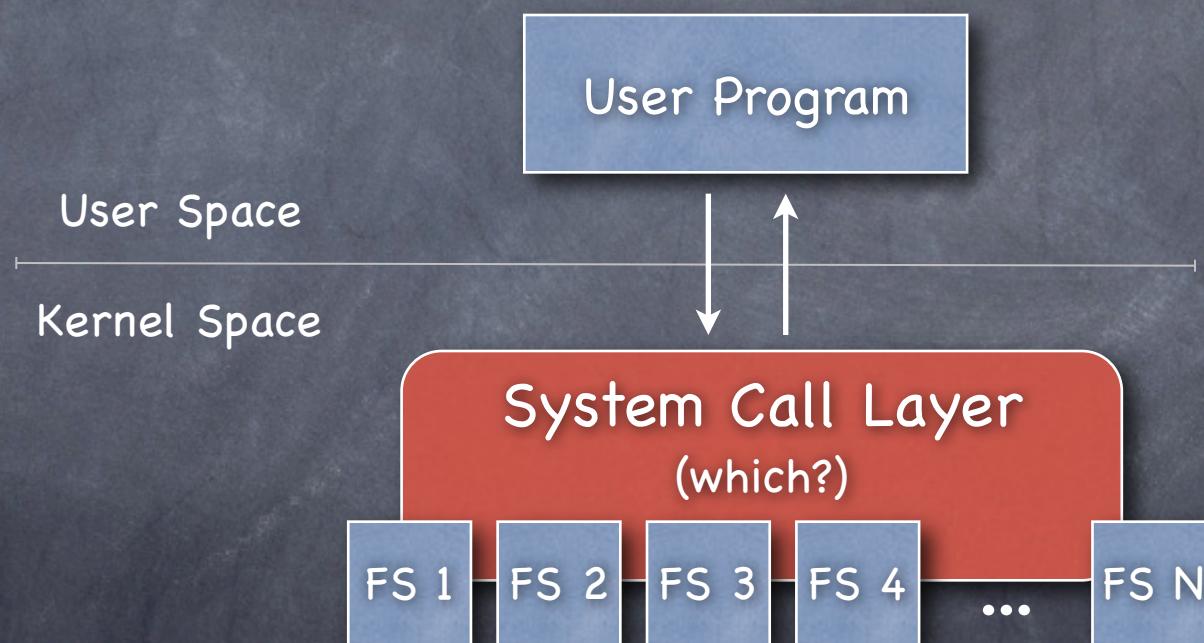
...A bit of History

- ⌚ Multics 1965 (File-System Paper)
A General-Purpose File System For Secondary Storage
- ⌚ Unix Late 1969



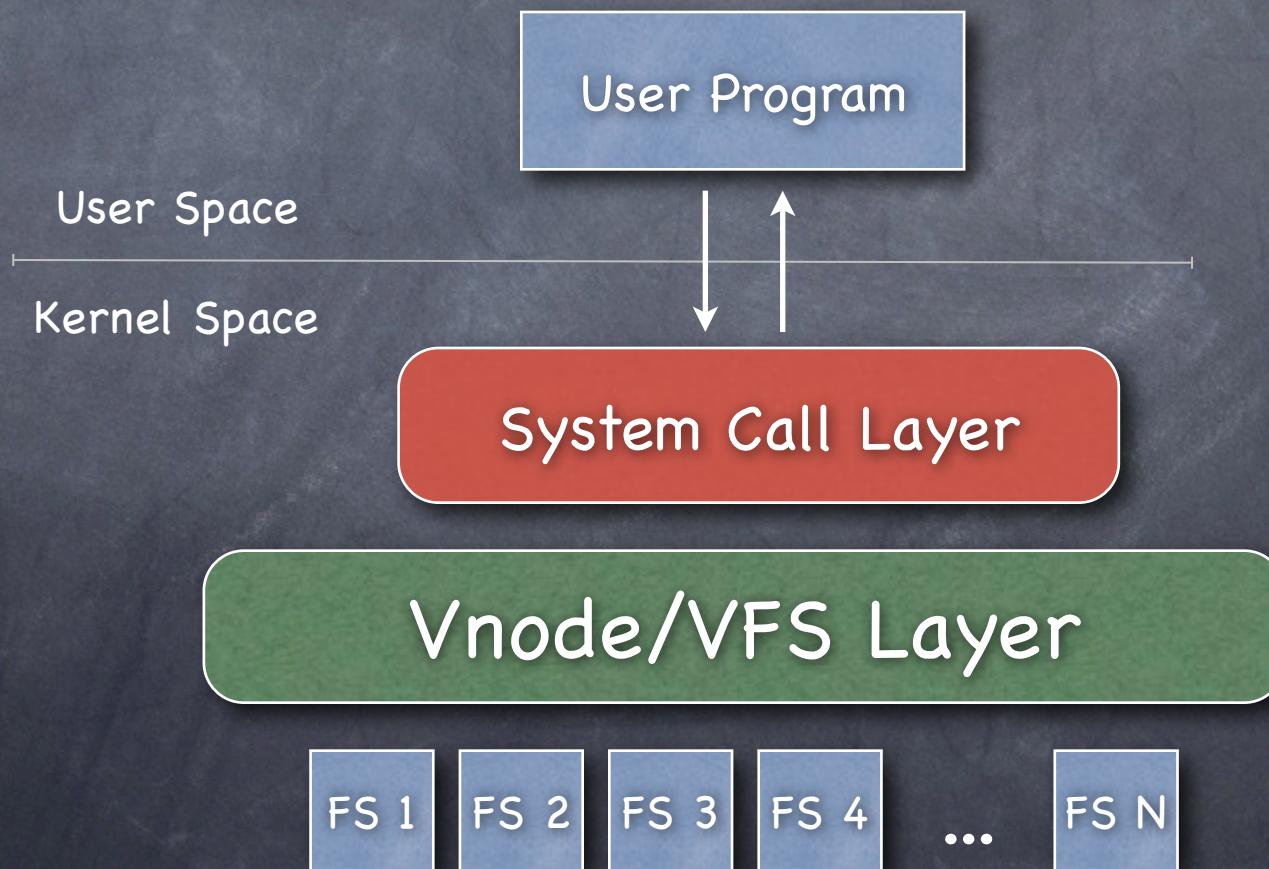
...A bit of History

- ⌚ Multics 1965 (File-System Paper)
A General-Purpose File System For Secondary Storage
- ⌚ Unix Late 1969



...A bit of History

- Multics 1965 (File-System Paper)
A General-Purpose File System For Secondary Storage
- Unix Late 1969
- Sun Microsystem 1984



Virtual File-System

- Provides an abstraction within the kernel which allows different filesystem implementations to coexist.
- Provides the filesystem interface to userspace programs.

User Space

C Library
(open(), read(), write(), ...)

Kernel Space

System Calls
(sys_open(), sys_read(), ...)

VFS Concepts

- A super-block object represents a filesystem.
- I-Nodes are filesystem objects such as regular files, directories, FIFOs, ...
- A file object represents a file opened by a process.

Kernel
Supported
File-Systems

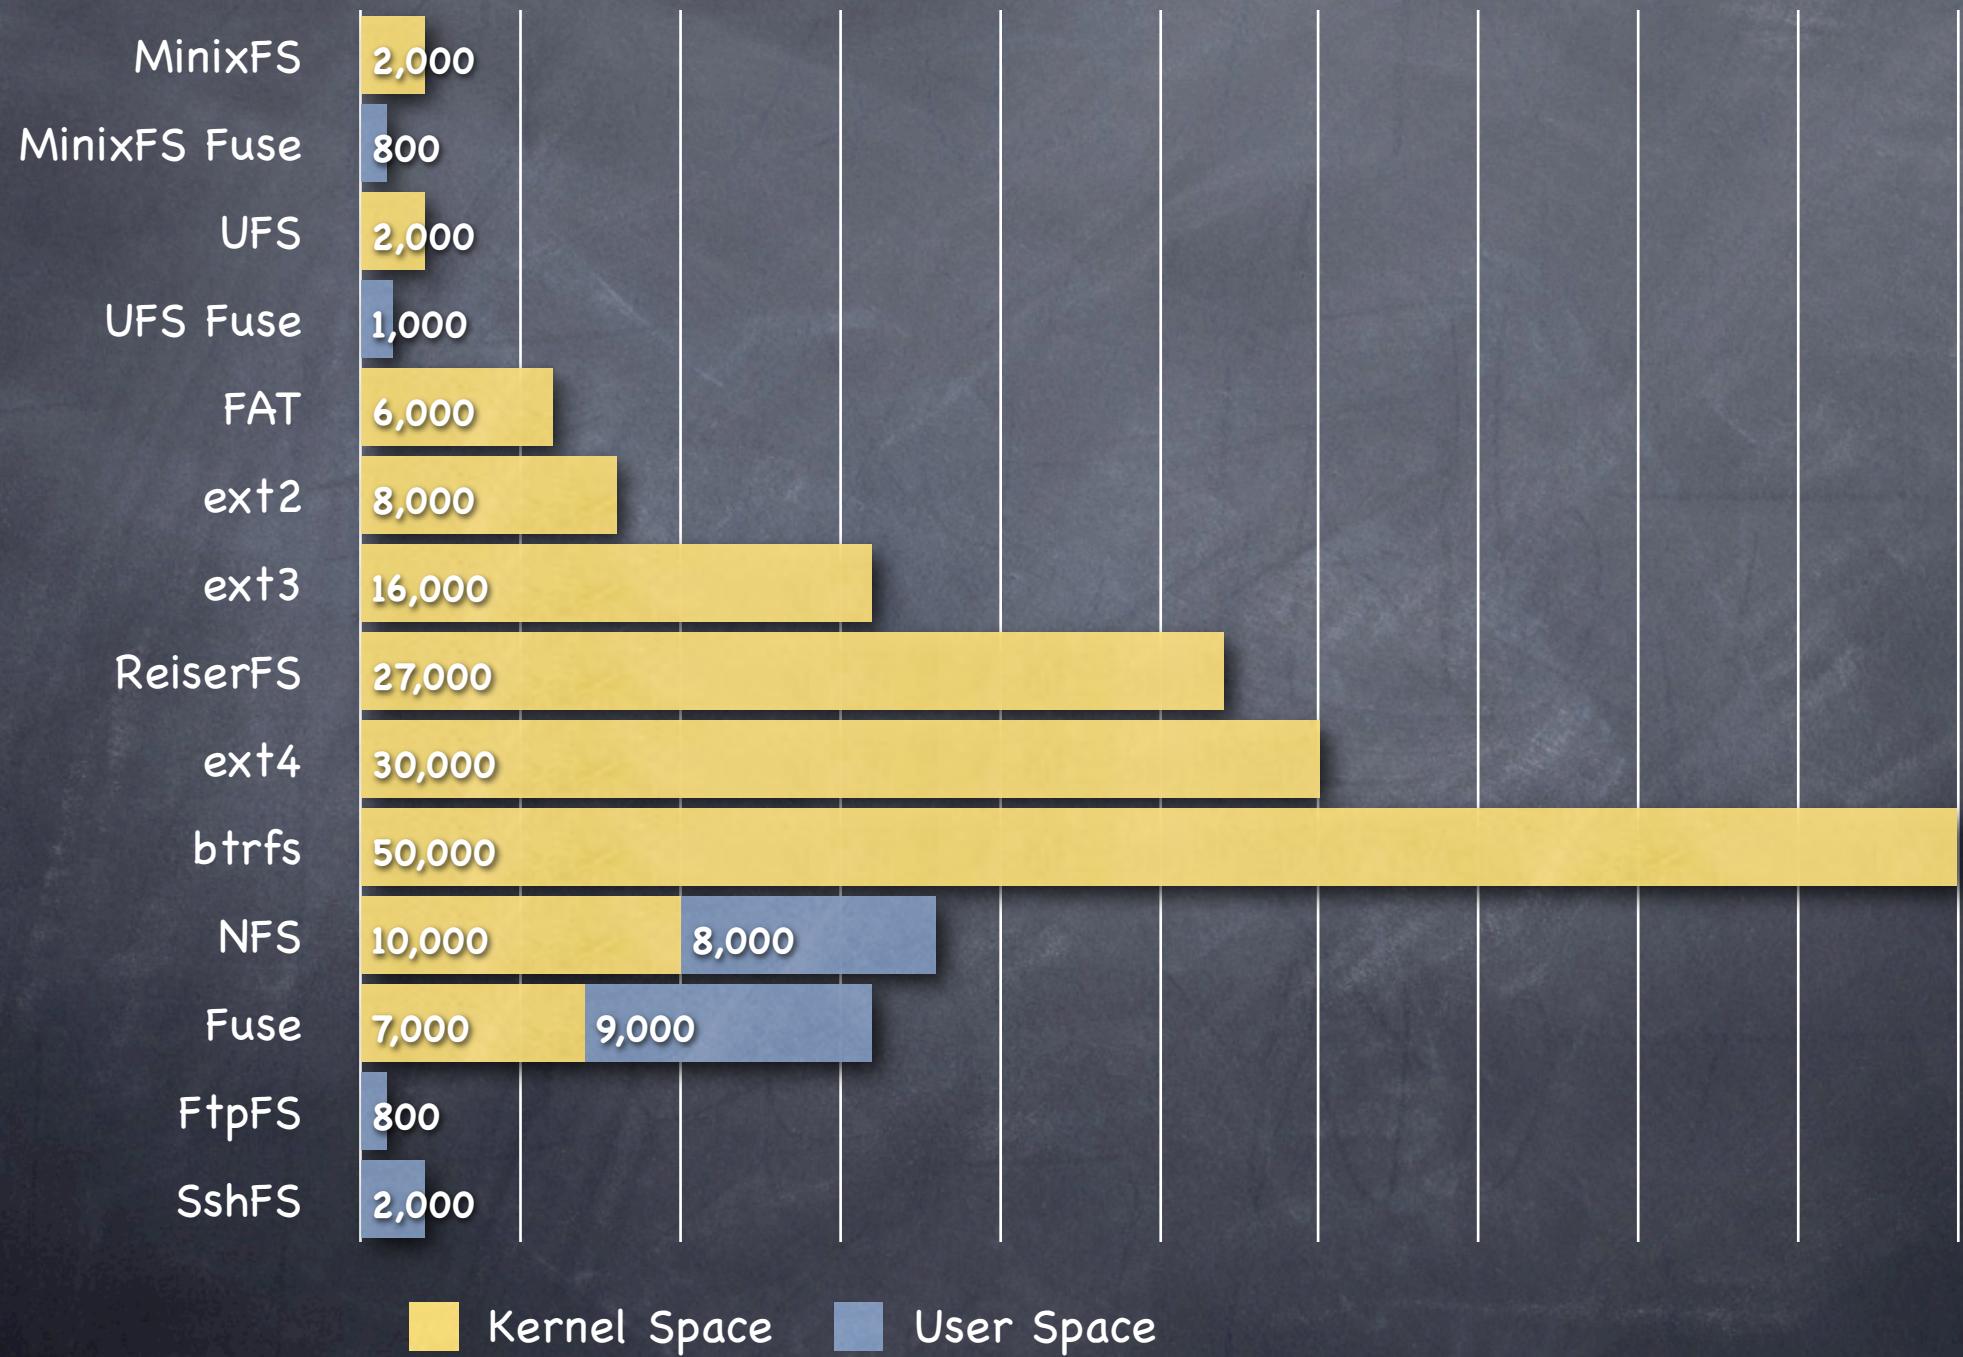
ext2	ReiserFS	XFS
ext3	Reiser4	JFS
ext4	Btrfs	HFS+
...

Wow, It seems not much
difficult writing a filesystem

Why File-System are Complex

- ⦿ You need to know the Kernel (No helper libraries: Qt, Glib, ...)
- ⦿ Reduce Disk Seek / SSD Block limited write cycles
- ⦿ Be consistent, Power Down, Unfinished Write...
(Journal, Soft-Updates, Copy-on-Write, ...)
- ⦿ Bad Blocks, Disk Error
- ⦿ Don't waste too much space for Metadata
- ⦿ Extra Features: Deduplication, Compression,
Cryptography, Snapshots...

File-Systems Lines of Code



Building a File-System is Difficult

- ⦿ Writing good code is not *easy* (Bugs, Typo, ...)
- ⦿ Writing good code in Kernel Space is much more difficult.
- ⦿ Too many reboots during the development
- ⦿ Too many Kernel Panic during Reboot
- ⦿ We need more flexibility and Speedups!

FUSE, develop your file-system
with your favorite language and library
in user space

What is FUSE

- ⦿ Kernel module! (like ext2, ReiserFS, XFS, ...)
- ⦿ Allows non-privileged user to create their own file-system without editing the kernel code. (User Space)
- ⦿ FUSE is particularly useful for writing "virtual file systems", that act as a view or translation of an existing file-system storage device. (Facilitate Disk-Based, Network-Based and Pseudo File-System)
- ⦿ Bindings: Python, Objective-C, Ruby, Java, C#, ...

File-Systems in User Space?

...Make File Systems Development Super Easy

- ⦿ All UserSpace Libraries are Available
- ⦿ ...Debugging Tools
- ⦿ No Kernel Recompilation
- ⦿ No Machine Reboot!
...File-System upgrade/fix
2 sec downtime, app restart!

Yeah, ...but what's FUSE?

It's a File-System with user-space callbacks

ntfs-3g

sshfs

ifuse

ChrionFS

gnome-vfs2

ftppfs

zfs-fuse

YouTubeFS

cryptoFS

RaleighFS
Unix

FUSE Kernel Space and User Space

The FUSE kernel module and the FUSE library communicate via a special file descriptor which is obtained by opening /dev/fuse



Beyond the Traditional File-Systems

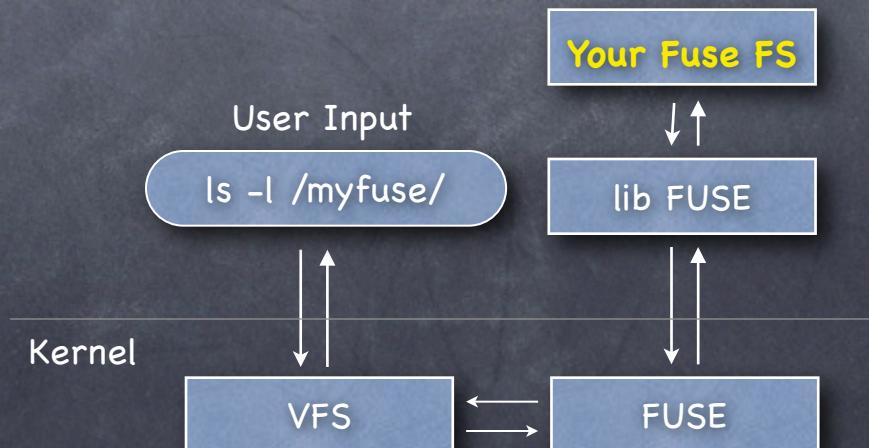
- ⦿ ImapFS: Access to your mail with grep.
- ⦿ SocialFS: Log all your social network to collect news/jokes and other social things.
- ⦿ YouTubeFS: Watch YouTube video as on your disk.
- ⦿ GMailFS: Use your mailbox as backup disk.

Thousand of tools available
cat/grep/sed

open() is the most used function in our applications

FUSE API Overview

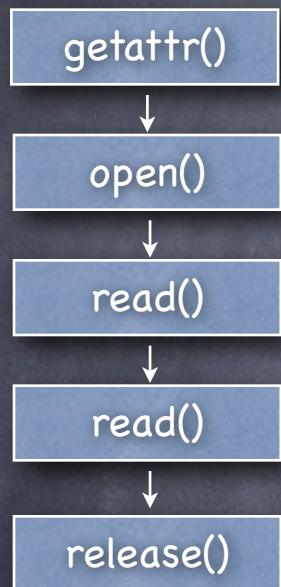
- ⦿ `create(path, mode)`
- ⦿ `truncate(path, size)`
- ⦿ `mknod(path, mode, dev)`
- ⦿ `open(path, mode)`
- ⦿ `write(path, data, offset)`
- ⦿ `read(path, length, offset)`
- ⦿ `release(path)`
- ⦿ `fsync(path)`
- ⦿ `chmod(path, mode)`
- ⦿ `chown(path, oid, gid)`
- ⦿ `mkdir(path, mode)`
- ⦿ `unlink(path)`
- ⦿ `readdir(path)`
- ⦿ `rmdir(path)`
- ⦿ `rename(opath, npath)`
- ⦿ `link(srcpath, dstpath)`



FUSE API Overview

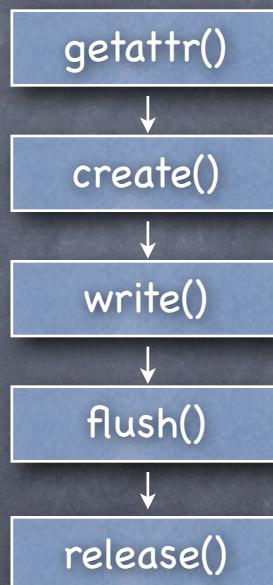
Reading

`cat /myfuse/test.txt`



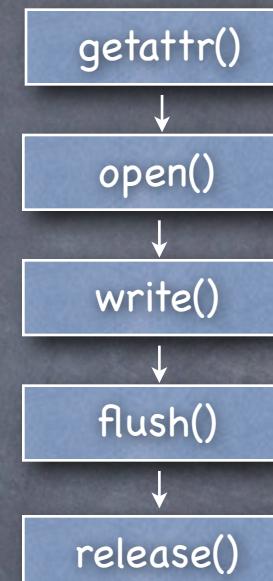
Writing

`echo Hello > /myfuse/test2.txt`



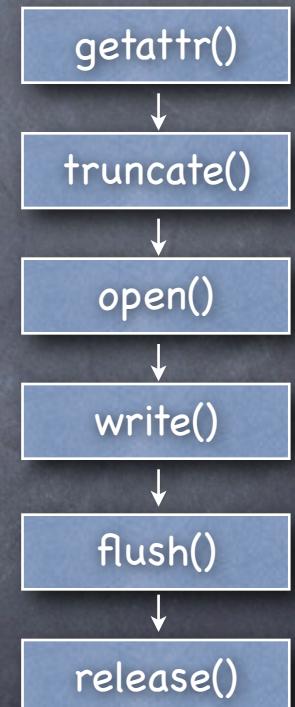
Appending

`echo World >> /myfuse/test2.txt`



Truncating

`echo Woo > /myfuse/test2.txt`



Removing

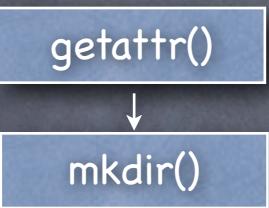
`rm /myfuse/test.txt`



FUSE API Overview

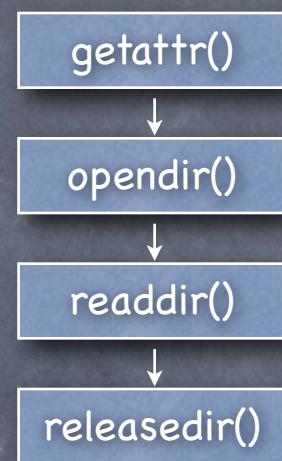
Creating

`mkdir /myfuse/folder`



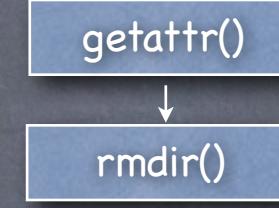
Reading

`ls /myfuse/folder/`



Removing

`rmdir /myfuse/folder`



Other Methods (`getattr()` is always called)

`chown th30z:develer /myfuse/test.txt` `getattr() -> chown()`
`chmod 755 /myfuse/test.txt` `getattr() -> chmod()`

`ln -s /myfuse/test.txt /myfuse/test-link.txt` `getattr() -> symlink()`

`mv /myfuse/folder /myfuse/fancy-folder` `getattr() -> rename()`

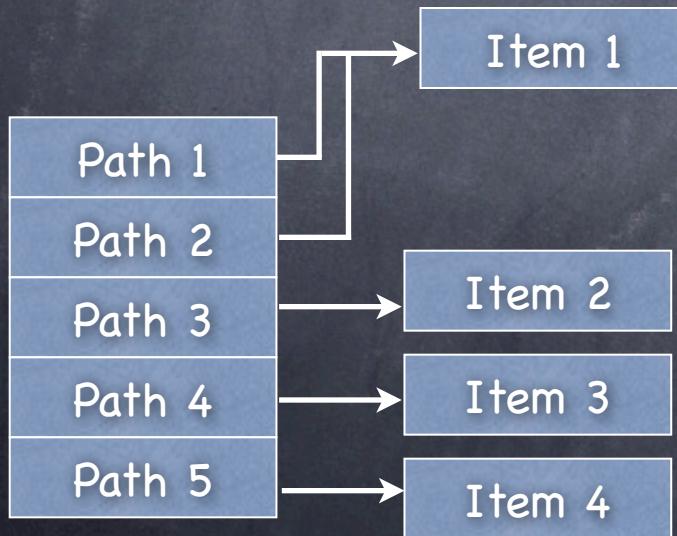
First Code Example!

HTFS

(HashTable File-System)

HTFS Overview

- Traditional Filesystem Object with Metadata (mode, uid, gid, ...)
- HashTable (dict) keys are paths values are Items.



(Disk - Storage HashTable)

FS Item/Object

Time of last access	Metadata
Time of last modification	
Time of last status change	
Protection and file-type (mode)	
User ID of owner (UID)	
Group ID of owner (GID)	
Extended Attributes (Key/Value)	

Data

- Item can be a Regular File or Directory or FIFO...
- Data is raw data or filename list if item is a directory.

HTFS Item

```
class Item(object):
    def __init__(self, mode, uid, gid):
        # ----- Metadata --
        self.atime = time.time()      # time of last acces
        self.mtime = self.atime       # time of last modification
        self.ctime = self.atime       # time of last status change

        self.mode = mode             # protection and file-type
        self.uid = uid               # user ID of owner
        self.gid = gid               # group ID of owner

    # Extended Attributes
    self.xattr = {}

    # --- Data -----
    if stat.S_ISDIR(mode):
        self.data = set()
    else:
        self.data = "
```

This is a File!
we've metadata
data and even xattr

HTFS Item

```
def read(self, offset, length):  
    return self.data[offset:offset+length]
```

```
def write(self, offset, data):  
    length = len(data)  
    self.data = self.data[:offset] + data + self.data[offset+length:]  
    return length
```

```
def truncate(self, length):  
    if len(self.data) > length:  
        self.data = self.data[:length]  
    else:  
        self.data += '\x00' * (length - len(self.data))
```

...a couple
of utility methods
to read/write
and interact with data.

HTFS Fuse Operations

```
class HTFS(fuse.Fuse):
    def __init__(self, *args, **kwargs):
        fuse.Fuse.__init__(self, *args, **kwargs)

        self.uid = os.getuid()
        self.gid = os.getgid()

        root_dir = Item(0755 | stat.S_IFDIR, self.uid, self.gid)
        self._storage = {'/': root_dir}
```

File-System must be initialized with
the / directory

```
def main():
    server = HTFS()
    server.main()
```

Your FUSE File-System
is like a Server...

getattr() is called before
any operation. Tells to
the VFS if you can access
to the specified file and
the "State".

```
def getattr(self, path):
    if not path in self._storage:
        return -errno.ENOENT
```

```
# Lookup Item and fill the stat struct
item = self._storage[path]
st = zstat(fuse.Stat())
st.st_mode = item.mode
st.st_uid = item.uid
st.st_gid = item.gid
st.st_atime = item.atime
st.st_mtime = item.mtime
st.st_ctime = item.ctime
st.st_size = len(item.data)
return st
```

HTFS Fuse Operations

```
def create(self, path, flags, mode):
    self._storage[path] = Item(mode | stat.S_IFREG, self.uid, self.gid)
    self._add_to_parent_dir(path)
```

```
def truncate(self, path, len):
    self._storage[path].truncate(len)
```

```
def read(self, path, size, offset):
    return self._storage[path].read(offset, size)
```

```
def write(self, path, buf, offset):
    return self._storage[path].write(offset, buf)
```

Disk is just a big dictionary...
...and files are items
key = name
value = data

```
def unlink(self, path):
    self._remove_from_parent_dir(path)
    del self._storage[path]
```

```
def rename(self, oldpath, newpath):
    item = self._storage.pop(oldpath)
    self._storage[newpath] = item
```

HTFS Fuse Operations

```
def mkdir(self, path, mode):
    self._storage[path] = Item(mode | stat.S_IFDIR, self.uid, self.gid)
    self._add_to_parent_dir(path)
```

```
def rmdir(self, path):
    self._remove_from_parent_dir(path)
    del self._storage[path]
```

```
def readdir(self, path, offset):
    dir_items = self._storage[path].data
    for item in dir_items:
        yield fuse.Direntry(item)
```

```
def _add_to_parent_dir(self, path):
    parent_path = os.path.dirname(path)
    filename = os.path.basename(path)
    self._storage[parent_path].data.add(filename)
```

Directory is a File
that contains
File names
as data!

HTFS Fuse Operations

```
def setxattr(self, path, name, value, flags):
    self._storage[path].xattr[name] = value
```

```
def getxattr(self, path, name, size):
    value = self._storage[path].xattr.get(name, "")
    if size == 0: # We are asked for size of the value
        return len(value)
    return value
```

```
def listxattr(self, path, size):
    attrs = self._storage[path].xattr.keys()
    if size == 0:
        return len(attrs) + len(".join(attrs)")
    return attrs
```

```
def removexattr(self, path, name):
    if name in self._storage[path].xattr:
        del self._storage[path].xattr[name]
```

Extended attributes extend
the basic attributes
associated with files and
directories in the file system.
They are stored as name:data
pairs associated with file
system objects

HTFS Fuse Operations

Lookup Item,
Access to its
information/data return
or write it.
This is the
File-System's Job

```
def symlink(self, path, newpath):
    item = Item(0644 | stat.S_IFLNK, self.uid, self.gid)
    item.data = path
    self._storage[newpath] = item
    self._add_to_parent_dir(newpath)

def readlink(self, path):
    return self._storage[path].data
```

```
def chmod(self, path, mode):
    item = self._storage[path]
    item.mode = mode

def chown(self, path, uid, gid):
    item = self._storage[path]
    item.uid = uid
    item.gid = gid
```

Symlinks contains just
pointed file path.

Other small Examples

Simulate Tera Byte Files

```
class TBFS(fuse.Fuse):
    def getattr(self, path):
        st = zstat(fuse.Stat())
        if path == '/':
            st.st_mode = 0644 | stat.S_IFDIR
            st.st_size = 1
            return st
        elif path == '/tera.data':
            st.st_mode = 0644 | stat.S_IFREG
            st.st_size = 128 * (2 ** 40)
            return st
        return -errno.ENOENT

    def read(self, path, size, offset):
        return '0' * size

    def readdir(self, path, offset):
        if path == '/':
            yield fuse.Dirent('tera.data')
```

Read-Only FS
with 1 file
of 128TiB

No
Disk/RAM Space
Required!

read()

Send data only
when is requested

X[^]OR File-System

```
def _xorData(data):
    data = [chr(ord(c) ^ 10) for c in data]
    return string.join(data, "")

class XorFS(fuse.Fuse):
    ...
    def write(self, path, buf, offset):
        data = _xorData(buf)
        return _writeData(path, offset, data)

    def read(self, path, length, offset):
        data = _readData(path, offset, length)
        return _xorData(data)

    ...
        res = _xorData("xor")
        print res // "rex"
        res2 = _xorData(res)
        print res2 // "xor"
```

Dup Write File-System

```
class DupFS(fuse.Fuse):
    def __init__(self, *args, **kwargs):
        ...
        fd_disk1 = open('/dev/hda1', ...)
        fd_disk2 = open('/dev/hdb5', ...)
        fd_log = open('/home/th30z/testfs.log', ...)
        fd_net = socket.socket(...)

        ...
    ...

    def write(self, path, buf, offset):
        ...
        disk_write(fd_disk1, path, offset, buf)
        disk_write(fd_disk2, path, offset, buf)
        net_write(fd_net, path, offset, buf)
        log_write(fd_log, path, offset, buf)

        ...
        ...do other fancy stuff
```

Write on your Disk partition 1 and 2.

Send data over Network

Log your file-system operations

One more thing

(File and Folders doesn't fit)

Rethink the File-System



I don't know
where I've to
place this file...

...Ok, for now
Desktop is a
good place...

(Mobile/Home Devices)

Rethink the File-System

Small Devices
Small Files
EMails, Text...

We need to
lookup quickly
our data. Tags,
Full-Text
Search...



...Encourage people
to view their content
as objects.

(Large Clusters, The Cloud...)

Rethink the File-System

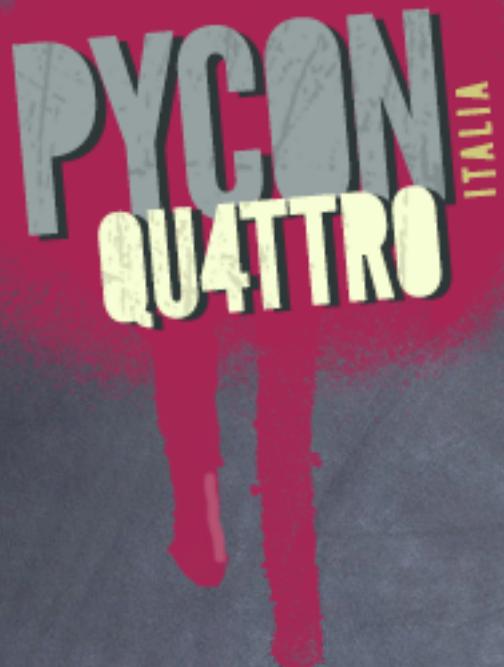
Distributed data

Scalability

Fail over

Cluster
Rebalancing





Q&A

Python FUSE

<http://mbertozzi.develer.com/python-fuse>